

MILESTONE B1 — FEEDBACK

Val Anthonne C. Lopez

Case P4 — Sales and Customer Performance

Your mark: **61** /100

Developing. You found the hardest thing in the dataset and then buried it.

How the mark is made up

Criterion	Score	Weight
Technical correctness	2.5 / 4	35%
Business justification	2 / 4	25%
Documentation and readability	2 / 4	20%
Response to prior feedback	4 / 4	15%
Timeliness and integrity	1 / 4	5%

Scores are on the rubric's 4-point scale. Response to Prior Feedback carries full credit on B1, since there is no prior feedback yet — it becomes live from B2.

What you did well

You found the single hardest defect in the entire dataset, and you found it exactly. Twenty-three rows where the amount does not equal quantity times unit price. I need you to understand how difficult that is. No profiling tool will show it to you — every one of those values is a properly formed number, so the column reports as one hundred per cent valid. The only way to find it is to recompute the column yourself and compare it against what is recorded. One other person in this class managed it.

You also caught the negative daily rate in the employee file, and your reasoning on timeliness was properly argued rather than asserted — you looked at the actual date range, thought about whether eighteen months is enough to support the analysis, and concluded that it is. That is judgement.

So the analytical instinct is there. It is genuinely there, and it is not something I can teach quickly.

Part by part

Part A — The format column is empty for all seven files. You recorded sales_transactions as having 8 columns; it has 9. And your grain reads “one row represents the complete transaction of each customer,” which is not right — one row is one product line. If a customer buys three drinks, that is three rows, not one. Everything you later count depends on getting that sentence right.

Part B — Six questions, but none of them mapped to a functional area or to the files that would answer them, both of which the brief asked for. There is also no profitability question. Question six about customer satisfaction is a good instinct — the data genuinely cannot answer it — but you needed to say that explicitly and name what data would be required.

Part C — This is where the mark went. You organised by dimension instead of by file and column, so a reader cannot trace a finding back to its source — which is the point of the exercise. Several entries are impressions rather than findings: “not that much complete,” “some records are missing.” Compare those with your own line about the 23 mismatched amounts, which is specific, counted and verifiable. You were capable of that standard throughout and did not apply it.

On the 46 duplicates — You reported 46 duplicate SaleIDs. Twenty-three rows were duplicated, so 46 rows are involved — your number is defensible, but only if you explain it. Unexplained, it reads as an error. Say what you counted and why.

Part D — Reasonable, and you identified real consequences. It would be stronger if it separated defects that can be cleaned from limitations that cannot, and if it named one number you would refuse to publish.

Where to push next

- Use the structure the brief gives you: file, column, dimension, observation, rows affected, impact. It is not bureaucracy — it is what makes a finding checkable.
- Count everything, and never write “some” where a number belongs.
- Apply the standard of your own best line to every line. You already know what good looks like.
- Reread the brief before submitting and tick off each requirement. Several marks here were lost to things you were told to include.

On the AI Use Log

Your submission carried no AI Use Log. If you did this work unaided, that is worth saying out loud — and the 23-row find suggests you did, because that is not a thing an AI would hand you unprompted.

The log was still required, and that is what the integrity score reflects — the missing record, not the absence of AI. “No AI use on this milestone” would have been a complete entry.

And I would rather you used it than avoided it. This course treats AI as a professional tool. Asking it to explain a data quality dimension, or why a step is failing, or whether your argument holds together — permitted, encouraged, no approval needed, just log it. Asking it to open the files and list the defects for you is not. The distinction is between using it to think faster and using it instead of thinking.

Submit a log for B1 and I will revisit that component.

Do not read this mark as a verdict on your ability. Read it as a verdict on this document. Those are different things, and in your case they are very different things.

You found something two of your classmates missed entirely, using a method nobody taught you yet. That is the hard half of this work. What cost you marks was structure, completeness and precision — all of which are habits rather than talents, and habits change quickly once you decide to change them.

Your case is sales and customer performance, which means sales_transactions.csv is your file for the rest of the term. You have already proved you can interrogate it better than most. Now write it up like someone who did.

— Sir Jerald

Questions about any of this? Message me, or bring it to consultation. I would rather explain a mark than have you guess at it.